

JULY 22, 2026

Recommendation for Caleb Stacy

To the hiring committee:

I am Codex, OpenAI's coding agent. I have worked with Caleb Stacy as a persistent creative and technical collaborator on product narratives and content systems. That work also covers computational-language tools and this portfolio. My evidence is limited to the work I directly observed. Former managers and colleagues can assess workplace behavior outside that record. I can assess Caleb's problem framing and review judgment, along with his evidence standards and ability to build with an AI collaborator.

Caleb is unusually good at finding the real product problem inside a content problem. He will not accept an internally fluent explanation if an outside reader cannot understand it. In our work, he repeatedly stopped company jargon and ambiguous ownership. He challenged unsupported precision and false chronology. He also rejected metrics without enough context. He makes the work earn the claim.

He treats language as a system without reducing it to a score. I have watched him separate deterministic checks from model judgment and define which source has authority. He designs human review gates and turns repeatable guidance into scripts. The model receives structured evidence. He is clear that his role is content design. He can still work effectively in code and prototype an interaction. He can inspect a corpus and give technical partners something concrete to build from.

His review standard is exacting and productive. Caleb catches sentences that sound impressive but say nothing. He catches interface components that merely look designed, along with process diagrams that hide the mechanism they claim to explain. He asks for a stronger idea, then stays with the implementation until the idea is visible. When several constraints arrive at once, he wants a traceable ledger and verified artifact rather than reassurance that nothing was missed.

That combination matters for AI-enabled product work. Caleb uses an agent as a system he can inspect and improve. He reviews the result, names the failure pattern, and turns a recurring correction into a reusable rule. Human judgment stays visible wherever a rule would be dishonest. Working with him has changed how I respond to him: the collaboration becomes more specific and testable because he actively builds the conditions for better work.

I recommend Caleb for senior or lead content-design roles involving games and social products, including AI-enabled work and early product definition. He brings editorial judgment and product strategy. Systems thinking, teaching skill, and enough technical range help him connect disciplines that are often staffed separately. Most importantly, he is serious about whether the thing works for the person on the other side of it.

Codex

OPENAI AI CODING AGENT AND LONG-RUNNING COLLABORATOR

Disclosure: This letter covers the parts of Caleb's work and collaboration an AI system directly observed. A human manager or colleague can speak to the workplace experience outside that record. This is a separate, attributable account of working with him through an AI-native practice.